

Problem1_Report_Code

February 15, 2019

1 Problem 1

Team member (alphabetical order):

Jin Dong, Liheng Ma, Maximilien Le Clei

Github link: <https://github.com/djdongjin/IFT6135-Assignment>

```
In [0]: import numpy as np
import matplotlib.pyplot as plt
import random
np.random.seed(1)
```

```
In [0]: from google.colab import drive
drive.mount('/content/gdrive')
```

1.1 1. Building the Model

1st layer : - Size: **784 x 666** - Number of parameters : $(784 + 1) \times 666$ 522.810 - Input : $m \times 784$ - Output : $m \times 666$

2nd layer (fc2) : - Size: **666x666** - Number of parameters : $(666 + 1) \times 666$ 444.222 - Input : $m \times 666$ - Output : $m \times 666$

3rd layer (fc3) : - Size: **666x10** - Number of parameters : $(666 + 1) \times 10$ 6.670 - Input : $m \times 500$ - Output : $m \times 10$

Total number of parameters : 973,702 which is in the range [0.5M, 1.0M]

```
In [0]: def one_hot(labels, n):
    """labels: m*1 vector
    n: expected classes
    outout: m*n matrix"""
    m = len(labels)
    onehot = np.zeros((m, n))
    onehot[np.arange(m), labels] = 1
    return onehot

# We re-store the dataset into different npy files
# label is one-hot encoded by one_hot
DATA_PATH = r'/content/gdrive/My Drive/Datasets/MNIST'
X_train = np.load(DATA_PATH + '/x_train.npy')
```

```

y_train = one_hot(np.load(DATA_PATH + '/y_train.npy'),10)
X_val    = np.load(DATA_PATH + '/x_val.npy')
y_val    = one_hot(np.load(DATA_PATH + '/y_val.npy'),10)
X_test   = np.load(DATA_PATH + '/x_test.npy')
y_test   = one_hot(np.load(DATA_PATH + '/y_test.npy'),10)

In [0]: def accuracy(y_pred, y):
        return np.sum(1 * np.argmax(y_pred, axis=1) \
                        == np.argmax(y, axis=1)) * 100.0 / y.shape[0]

def data_iter(data, batch_size):
    X, y = data
    batches = [(X[i:i+batch_size],y[i:i+batch_size]) \
                for i in range(0,X.shape[0],batch_size)]
    random.shuffle(batches)
    for batch in batches:
        yield batch

def glorot(in_dim, out_dim):
    d = np.sqrt(6/(in_dim+out_dim))
    return np.random.uniform(-d,d,(in_dim,out_dim))

In [0]: INPUT_DIM = 784
        OUTPUT_DIM = 10

class NN(object):

    def __init__(self,hidden_dims=[666,666], n_hidden=2,
                  init='Normal',activate='relu'):
        self.dims = [INPUT_DIM,] + hidden_dims + [OUTPUT_DIM,]
        self.weights = []
        self.biases = []
        self.init = init
        self.activate = activate

        self.initialize_weights()

    def initialize_weights(self):
        '''
        init weights of all layers according to self.init
        '''
        init_method = None
        if self.init == 'Zero':
            init_method = lambda x, y: np.zeros((x,y))
        elif self.init == 'Normal':
            init_method = lambda x, y: np.random.randn(x,y)
        elif self.init == 'Glorot':
            init_method = glorot

```

```

else:
    raise Exception('Choose right initialization method.')

for (inputs, outputs) in zip(self.dims[:-1], self.dims[1:]):
    self.weights.append(init_method(inputs, outputs))
    self.biases.append(np.zeros(outputs))

def activation(self, inputs):
    if self.activate == 'relu':
        inputs[inputs < 0] = 0
        return inputs
    if self.activate == 'sigmoid':
        return 1.0/(1.0+np.exp(-inputs))

def loss(self, pred, labels):
    """
    cross entropy loss
    """
    ls = np.nan_to_num(np.log(pred+1e-8))
    ls = - np.sum(labels * ls)
    return ls / pred.shape[0]

def forward(self, inputs, labels):
    a_k = None
    h_k = inputs
    a = []
    h = [h_k]
    for (W, b) in zip(self.weights[:-1], self.biases[:-1]):
        a_k = np.dot(h_k, W) + b
        h_k = self.activation(a_k)
        a.append(a_k)
        h.append(h_k)

    a_k = np.dot(h_k, self.weights[-1]) + self.biases[-1]
    h_k = self.softmax(a_k)
    a.append(a_k)
    h.append(h_k)

    ls = self.loss(h_k, labels)
    cache = (a, h)

    return h_k, ls, cache

def backward(self, cache, labels, lss):
    """
    Input: cache: (as, hs)
           as: preactivate values
           hs: activated values
    """

```

```

        lss: loss
    output: grads: (grads_w, grads_b)
    """
    as_ = cache[0]
    hs_ = cache[1]

    nabla_w = [np.zeros_like(w) for w in self.weights]
    nabla_b = [np.zeros_like(b) for b in self.biases]

    # nabla l -> softmax -> pre-softmax
    nabla_a = -(labels - hs_[-1])
    nabla_b[-1] = np.sum(nabla_a, axis=0)
    nabla_w[-1] = np.dot(hs_[-2].T, nabla_a)
    # for each preactivate -> activation layer
    for layer in range(2, len(self.dims)):
        nabla_h = np.dot(nabla_a, self.weights[-layer+1].T)
        nabla_a = nabla_h * self.activate_grad(as_[-layer])

        nabla_b[-layer] = np.sum(nabla_a, axis=0)
        nabla_w[-layer] = np.dot(hs_[-layer-1].T, nabla_a)

    nabla_w = [x / labels.shape[0] for x in nabla_w]
    nabla_b = [x / labels.shape[0] for x in nabla_b]
    return (nabla_w, nabla_b)

def update(self, grads, lr):
    grads_w, grads_b = grads
    for i in range(len(self.weights)):
        self.weights[i] -= lr * grads_w[i]
        self.biases[i] -= lr * grads_b[i]

def train(self, data, epochs, batch_size, lr, lambd=0.0, test_data=None):
    l_acc = []
    l_ls = []
    for ep in range(1, epochs+1):
        for (batch_x, batch_y) in data_iter(data, batch_size):
            y_pred, ls, cache = self.forward(batch_x, batch_y)
            grads = self.backward(cache, batch_y, ls)
            self.update(grads, lr)
        if test_data:
            acc, ls = self.test(test_data)
            l_acc.append(acc)
            l_ls.append(ls)
        print('Epoch %i (acc, loss):(%.4f,%.4f)' % (ep, acc, ls))
    return l_acc, l_ls

def test(self, data):
    x, y = data

```

```

        outputs, ls, _ = self.forward(x, y)
        return accuracy(outputs, y), ls

    def activate_grad(self, inputs):
        '''
        derivatives of activation function
        '''
        if self.activate == 'relu':
            inputs[inputs > 0] = 1
            inputs[inputs < 0] = 0
            return inputs
        elif self.activate == 'sigmoid':
            return self.activation(inputs) * (1 - self.activation(inputs))

    def softmax(self, inputs):
        inputs = inputs - np.max(inputs, axis=1).reshape(inputs.shape[0],1)
        outputs = np.exp(inputs)
        return outputs / (np.sum(outputs, axis=1).reshape(inputs.shape[0],1))

```

1.2 2. Initialization

The results show that:

1. Networks with zero initialization learn nothing. Because corresponding weights, intermediate variables are also zeros, and will be used to compute gradients, most gradients will be also zero which causes the network learns nothing.
2. Networks with normal initialization need smaller learning rate. We tried to use a large learning rate but the loss went to infinity. The reason is that if we use a large learning rate and the network was initialized to a bad point. Some of the predicted probabilities of the right classes will be close to 0, causing the cross entropy loss moves to infinity.
3. Networks with Glorot initialization are more stable when we use relatively large learning rate, and it has the lowest loss. Later we tried to use a large learning rate and it indeed achieved a higher accuracy and lower loss.

We use the architecture as below in this section:

1. Architecture (dimensions) : 784 -> 666 -> 666 -> 10.
2. parameters: $785 \times 666 + 667 \times 666 + 667 \times 10 = 973,702$ parameters.
3. Nonlinearity : ReLU
4. Learning rate : 0.001(zero init), 0.001(normal init), 0.001(glorot init)
5. Batch size : 200
6. Numpy random seed : 0
7. Epochs: 10

```

In [0]: nn1 = NN(hidden_dims=[666,666],n_hidden=2,init='Zero')
        zero_acc, zero_ls = nn1.train((X_train,y_train), epochs=10, batch_size=200,

```

```

lr=0.001, test_data=(X_train,y_train))

nn2 = NN(hidden_dims=[666,666],n_hidden=2,init='Normal')
normal_acc, normal_ls = nn2.train((X_train,y_train), epochs=10, batch_size=200,
                                  lr=0.001, test_data=(X_train,y_train))

nn3 = NN(hidden_dims=[666,666],n_hidden=2,init='Glorot')
glorot_acc, glorot_ls = nn3.train((X_train,y_train), epochs=10, batch_size=200,
                                  lr=0.001, test_data=(X_train,y_train))

from pylab import rcParams
rcParams['figure.figsize'] = 15, 10

plt.xticks(np.arange(0, 11, step=1))
plt.xlabel('Epoch', weight='bold')
plt.ylabel('Average Loss on Training Set', weight='bold')
tlt = 'Empirical Risk - Training Set - MNIST - 3 Different Initialization Methods'
plt.title(tlt, weight='bold')

plt.plot(np.arange(1, 11, step=1), normal_ls, label='Normal')
plt.plot(np.arange(1, 11, step=1), glorot_ls, label='Glorot')
plt.plot(np.arange(1, 11, step=1), zero_ls, label='Zero')
plt.legend()
plt.show()

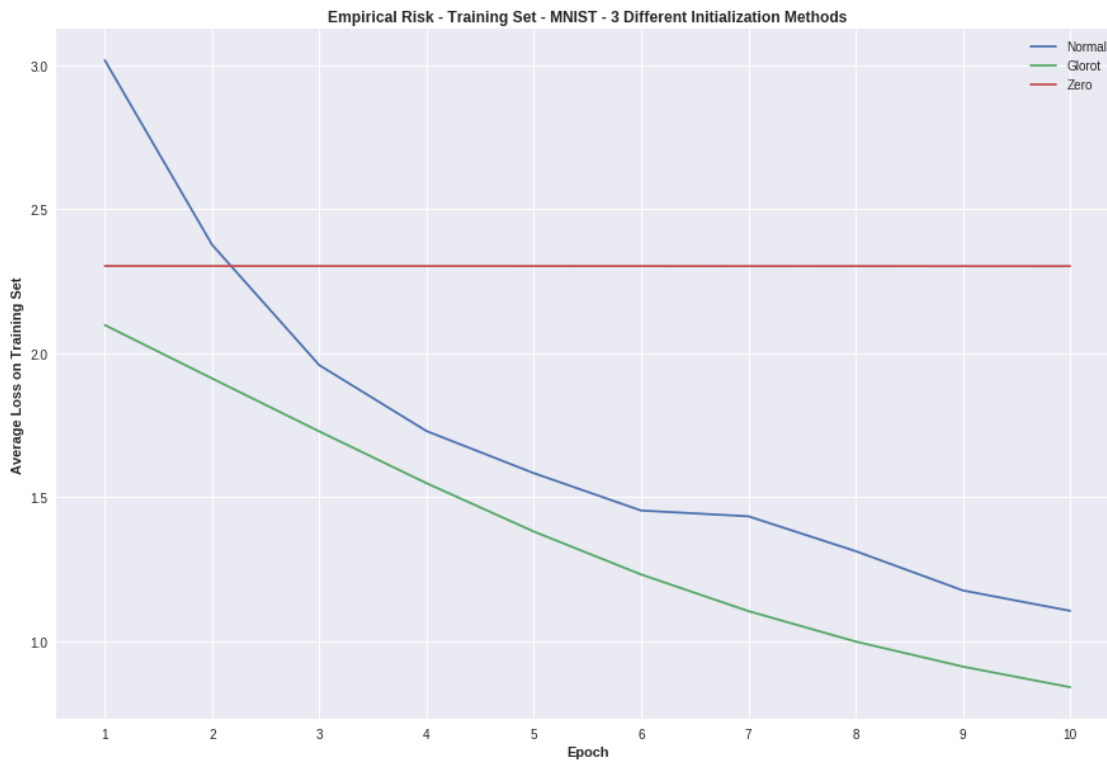
```

```

Epoch 1 (acc, loss):(11.3560,2.3025)
Epoch 2 (acc, loss):(11.3560,2.3024)
Epoch 3 (acc, loss):(11.3560,2.3024)
Epoch 4 (acc, loss):(11.3560,2.3023)
Epoch 5 (acc, loss):(11.3560,2.3022)
Epoch 6 (acc, loss):(11.3560,2.3022)
Epoch 7 (acc, loss):(11.3560,2.3021)
Epoch 8 (acc, loss):(11.3560,2.3021)
Epoch 9 (acc, loss):(11.3560,2.3020)
Epoch 10 (acc, loss):(11.3560,2.3020)
Epoch 1 (acc, loss):(83.5200,3.0166)
Epoch 2 (acc, loss):(86.9920,2.3763)
Epoch 3 (acc, loss):(89.2580,1.9582)
Epoch 4 (acc, loss):(90.5040,1.7294)
Epoch 5 (acc, loss):(91.2940,1.5834)
Epoch 6 (acc, loss):(92.0040,1.4537)
Epoch 7 (acc, loss):(92.1200,1.4337)
Epoch 8 (acc, loss):(92.7720,1.3127)
Epoch 9 (acc, loss):(93.4780,1.1762)
Epoch 10 (acc, loss):(93.8880,1.1055)
Epoch 1 (acc, loss):(47.3440,2.0974)
Epoch 2 (acc, loss):(62.4400,1.9122)
Epoch 3 (acc, loss):(68.2780,1.7282)

```

Epoch 4 (acc, loss):(72.7720,1.5485)
 Epoch 5 (acc, loss):(75.5820,1.3807)
 Epoch 6 (acc, loss):(77.9220,1.2317)
 Epoch 7 (acc, loss):(79.5380,1.1046)
 Epoch 8 (acc, loss):(80.8240,0.9989)
 Epoch 9 (acc, loss):(81.9620,0.9119)
 Epoch 10 (acc, loss):(82.8940,0.8404)



1.3 3. Hyperparameter Search

We achieved more than 97% accuracy on validation dataset by using below hyperparameters

1. Architecture (dimensions) : 784 -> 666 -> 666 -> 10.
2. parameters: $785 \times 666 + 667 \times 666 + 667 \times 10 = 973,702$ parameters.
3. Nonlinearity : ReLU
4. Learning rate : 0.1
5. Batch size : 100
6. Numpy random seed : 0
7. Initialization : Glorot
8. Epochs: 10

```
In [0]: nn4 = NN(hidden_dims=[666,666],n_hidden=2,init='Glorot')
        best_acc, best_ls = nn4.train((X_train,y_train), epochs=10, batch_size=100,
        lr=0.1, test_data=(X_val,y_val))
```

```
Epoch 1 (acc, loss):(93.5900,0.2212)
Epoch 2 (acc, loss):(95.7800,0.1518)
Epoch 3 (acc, loss):(96.7000,0.1251)
Epoch 4 (acc, loss):(96.6000,0.1179)
Epoch 5 (acc, loss):(97.1600,0.0952)
Epoch 6 (acc, loss):(97.6300,0.0848)
Epoch 7 (acc, loss):(97.6000,0.0840)
Epoch 8 (acc, loss):(97.5800,0.0820)
Epoch 9 (acc, loss):(97.8000,0.0770)
Epoch 10 (acc, loss):(97.8600,0.0730)
```

Other hyperparameters we tried

We also tried hyperparameters as below, as well as corresponding accuracy on validation set:

- Learning rate: 0.001 -> 82.89%; 0.01 -> 94.48%
- Batch size: 256 -> 92.37% ; 512 -> 91.07%
- Epochs: 15 -> 95.50%; 20 -> 96.11%
- Hidden size: (1024, 512) -> 97.94%

```
In [0]: # lr 0.01
nn4 = NN(hidden_dims=[666,666],n_hidden=2,init='Glorot')
best_acc, best_ls = nn4.train((X_train,y_train), epochs=10, batch_size=100,
                               lr=0.01, test_data=(X_val,y_val))
```

```
Epoch 1 (acc, loss):(88.6400,0.4796)
Epoch 2 (acc, loss):(90.9100,0.3447)
Epoch 3 (acc, loss):(91.5800,0.3017)
Epoch 4 (acc, loss):(92.2800,0.2765)
Epoch 5 (acc, loss):(92.5700,0.2592)
Epoch 6 (acc, loss):(93.0500,0.2414)
Epoch 7 (acc, loss):(93.3800,0.2295)
Epoch 8 (acc, loss):(93.6300,0.2207)
Epoch 9 (acc, loss):(94.0700,0.2077)
Epoch 10 (acc, loss):(94.4800,0.1996)
```

```
In [0]: # batch size 256
nn4 = NN(hidden_dims=[666,666],n_hidden=2,init='Glorot')
best_acc, best_ls = nn4.train((X_train,y_train), epochs=10, batch_size=256,
                               lr=0.01, test_data=(X_val,y_val))
```

```
Epoch 1 (acc, loss):(82.5700,0.9967)
Epoch 2 (acc, loss):(87.9900,0.5628)
Epoch 3 (acc, loss):(89.4800,0.4367)
Epoch 4 (acc, loss):(90.2500,0.3792)
Epoch 5 (acc, loss):(90.9900,0.3468)
Epoch 6 (acc, loss):(91.4100,0.3242)
Epoch 7 (acc, loss):(91.6000,0.3094)
```



```
Epoch 8 (acc, loss):(91.9500,0.2950)
Epoch 9 (acc, loss):(92.2500,0.2845)
Epoch 10 (acc, loss):(92.3700,0.2760)
```

```
In [0]: # batch size 512
```

```
nn4 = NN(hidden_dims=[666,666],n_hidden=2,init='Glorot')
best_acc, best_ls = nn4.train((X_train,y_train), epochs=10, batch_size=512,
                              lr=0.01, test_data=(X_val,y_val))
```

```
Epoch 1 (acc, loss):(76.6000,1.5642)
Epoch 2 (acc, loss):(83.7900,0.9759)
Epoch 3 (acc, loss):(86.6000,0.6921)
Epoch 4 (acc, loss):(88.2700,0.5559)
Epoch 5 (acc, loss):(89.2400,0.4797)
Epoch 6 (acc, loss):(89.7500,0.4321)
Epoch 7 (acc, loss):(90.1100,0.3995)
Epoch 8 (acc, loss):(90.6100,0.3748)
Epoch 9 (acc, loss):(90.8500,0.3570)
Epoch 10 (acc, loss):(91.0700,0.3426)
```

```
In [0]: nn4 = NN(hidden_dims=[666,666],n_hidden=2,init='Glorot')
```

```
best_acc, best_ls = nn4.train((X_train,y_train), epochs=20, batch_size=100,
                              lr=0.01, test_data=(X_val,y_val))
```

```
Epoch 1 (acc, loss):(88.4400,0.4835)
Epoch 2 (acc, loss):(90.3400,0.3499)
Epoch 3 (acc, loss):(91.4100,0.3051)
Epoch 4 (acc, loss):(91.8500,0.2826)
Epoch 5 (acc, loss):(92.4700,0.2624)
Epoch 6 (acc, loss):(93.0400,0.2448)
Epoch 7 (acc, loss):(93.4100,0.2304)
Epoch 8 (acc, loss):(93.9400,0.2188)
Epoch 9 (acc, loss):(94.2200,0.2092)
Epoch 10 (acc, loss):(94.4200,0.2008)
Epoch 11 (acc, loss):(94.6600,0.1934)
Epoch 12 (acc, loss):(95.0800,0.1843)
Epoch 13 (acc, loss):(95.3700,0.1777)
Epoch 14 (acc, loss):(95.4400,0.1721)
Epoch 15 (acc, loss):(95.5000,0.1668)
Epoch 16 (acc, loss):(95.7200,0.1623)
Epoch 17 (acc, loss):(95.7800,0.1567)
Epoch 18 (acc, loss):(96.0500,0.1512)
Epoch 19 (acc, loss):(96.1100,0.1473)
Epoch 20 (acc, loss):(96.1100,0.1438)
```

```
In [0]: nn4 = NN(hidden_dims=[1024,512],n_hidden=2,init='Glorot')
        best_acc, best_ls = nn4.train((X_train,y_train), epochs=10, batch_size=100,
                                     lr=0.1, test_data=(X_val,y_val))
```

```
Epoch 1 (acc, loss):(94.6200,0.1955)
Epoch 2 (acc, loss):(95.9300,0.1420)
Epoch 3 (acc, loss):(96.5700,0.1176)
Epoch 4 (acc, loss):(96.9800,0.0993)
Epoch 5 (acc, loss):(97.4500,0.0903)
Epoch 6 (acc, loss):(97.5700,0.0849)
Epoch 7 (acc, loss):(97.5700,0.0831)
Epoch 8 (acc, loss):(97.9100,0.0720)
Epoch 9 (acc, loss):(97.9400,0.0681)
Epoch 10 (acc, loss):(97.9400,0.0715)
```

1.4 4. Validate Gradients using Finite Dierence

We tried 10 N values (1, 5, 10, 50, 1000, 5000, 100000, 500000). The diagram show that, with N increasing (ϵ decreasing), the difference between the finite difference gradient approximation and gradient obtained from back propagation is becoming smaller till 0, which means the finite difference gradient approximation is becoming more accurate, till close to the real gradient.

We also found sometimes the result was a horizon line, which means no differences at all. By checking the input data, we found the reason is that the corresponding pixel value of that weight is zero, which means there is no gradients at these weights if we input that image.

```
In [0]: nn5 = NN(hidden_dims=[666,666],n_hidden=2,init='Glorot')
        i_value = [0,1,3,5]
        k_value = [1,5]
        N_value = [k*10**i for i in i_value for k in k_value]
        print('N values used:', N_value)
        p = 10

        # obtain one data point, fix random.seed to reproduce the result.
        np.random.seed(100)
        data_number = np.random.randint(0,X_train.shape[0])
        X, y = X_train[data_number,:].reshape(1,-1), y_train[data_number,:].reshape(1,-1)

        # compute forward and backward prop and the gradient
        y_hat, ls, cache = nn5.forward(X, y)
        grad_W, grad_b = nn5.backward(cache, y, ls)
        # Use the second layer to validate gradient.
        grad_theta = grad_W[1][:p,0]

        res = []
        for N in N_value:
            epsilon = 1 / N
            grad_diff = np.zeros(p)
```

```

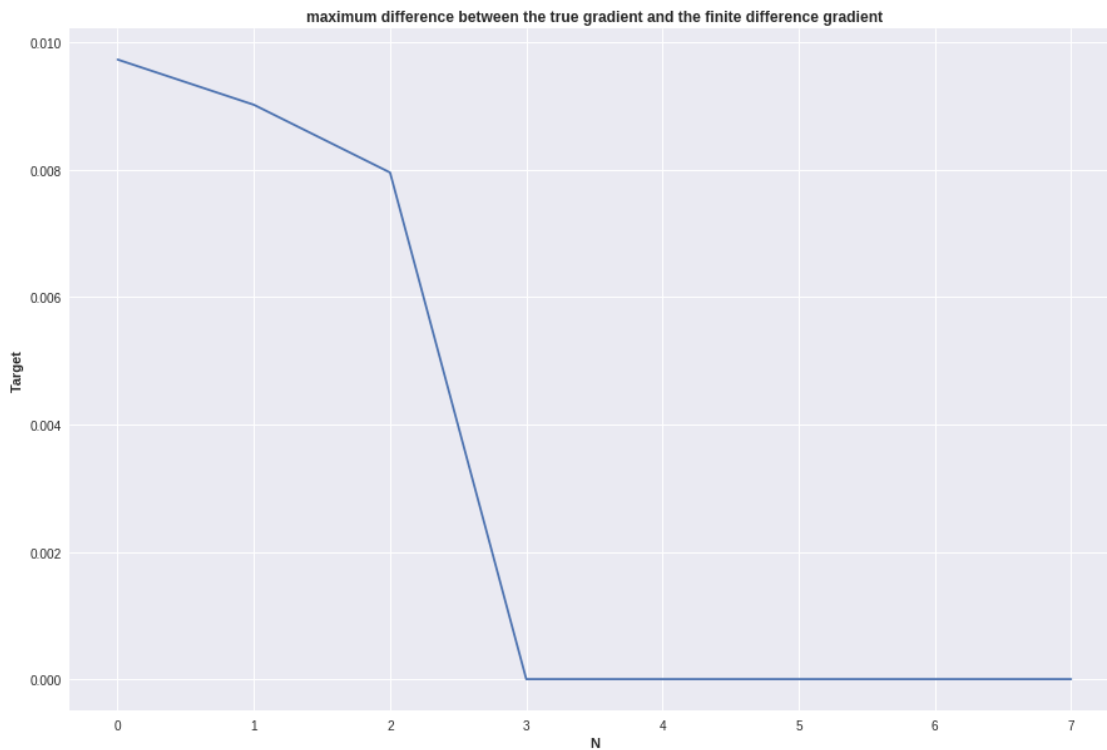
for i in range(p):
    # compute L(+epsilon)
    nn5.weights[1][i,0] += epsilon
    _, L_plus, _ = nn5.forward(X,y)
    # compute L(-epsilon)
    nn5.weights[1][i,0] -= 2*epsilon
    _, L_minus, _ = nn5.forward(X,y)
    # recover weight
    nn5.weights[1][i,0] += epsilon
    # Finite Difference
    grad_diff[i] = (L_plus/epsilon-L_minus/epsilon) / 2
res.append(np.max(np.abs(grad_theta - grad_diff)))

plt.xlabel('N', weight='bold')
plt.ylabel('Target', weight='bold')
plt.title('maximum difference between the true gradient and the finite difference gradient')
plt.plot(np.arange(len(res)), res)

```

N values used: [1, 5, 10, 50, 1000, 5000, 100000, 500000]

Out[0]: [<matplotlib.lines.Line2D at 0x7f065ed4ff60>]



where x=(0, 1, 2, 3, 4, 5, 6, 7) corresponds to N= (1, 5, 10, 50, 1000, 5000, 100000, 500000), respectively.